

# Data-Driven Discovery of Mechanistic Ecosystem Models with LLMs

October 17, 2025

**Keywords:** Artificial Intelligence, Ecological modelling, Evolutionary Algorithms, Large Language Models, Marine Ecosystems, Crown-of-Thorns Starfish, Great Barrier Reef, time-series forecasting

## Abstract

1. Ecosystem models are essential for ecosystem management, but their development traditionally requires significant time and expertise, creating bottlenecks in addressing urgent environmental challenges.
2. We present ***LEMMA*** (LLM Enabled Mechanistic Modelling for ecosystem Assessment), a framework that programmatically generates and iteratively refines mechanistic ecosystem models by combining large language models (LLMs) for equation synthesis and parameter search, evolutionary algorithms for structural optimization, and Template Model Builder (TMB) for efficient parameter estimation.
3. We critically review *LEMMA*'s ability to recover known ecological relationships through two complementary marine case studies: (1) a nutrient-phytoplankton-zooplankton model, and (2) a Crown-of-Thorns starfish (COTS) model. In the first case, our best models displayed almost perfect recovery of known ecological dynamics while maintaining strong predictive performance across multivariate time-series. In the second case, best *LEMMA* generated models approached human expert models in terms of their ability to successfully capture COTS outbreak dynamics and demonstrated strong out-of-sample predictive power.
4. *LEMMA* produces interpretable models with meaningful parameters that capture real biological processes, facilitating scientific insight and potentially accelerating management applications. By dramatically accelerating model development while offering ecological interpretability, *LEMMA* offers a powerful new tool for addressing urgent ecological challenges in a changing world.

## Data/Code Availability

The data and code associated with this manuscript is available at the anonymised repository: <https://anonymous.4open.science/r/EMs-with-LLMs-10E8>

# 1 Introduction

Ecosystem models provide invaluable information for managing complex interactions between nature and people (McCarthy and Possingham, 2004; Holden and Ellner, 2016), but their development traditionally requires significant time and expertise, creating a bottleneck in addressing urgent environmental challenges (Dichmont et al., 2017; Holden et al., 2024), particularly as climate change demands rapid, adaptable approaches for ecosystem management (Weiskopf et al., 2020; Malhi et al., 2020).

Artificial Intelligence (AI) offers great promise as a part of the solution to these modelling challenges, with potential to accelerate model development and enhance adaptability (Spillias et al., 2024a). While initial efforts to apply AI in ecological modelling focused on machine learning approaches that rely on black-box methods (Morales-García et al., 2024), emerging techniques in equation discovery and automated scientific discovery show particular promise (Huntingford et al., 2024; Floryan and Graham, 2022). These methods can derive interpretable mathematical relationships directly from data, offering advantages over statistical emulators when modelling novel environmental conditions (Schaeffer, 2017; Chen et al., 2024; Karniadakis et al., 2021).

A recent development in AI is the rise of large language models (LLMs): neural networks trained on vast text corpora, including scientific literature, code, and web content. LLMs can generate and interpret language, code, and equations, encoding broad domain knowledge that can be leveraged for scientific modelling tasks (Mammides and Papadopoulos, 2024; Wills et al., 2024). Attempts to leverage LLMs for direct time-series prediction (Zhang et al., 2024; Su et al., 2024; Hassani and Silva, 2024; Gandhi et al., 2024; Bylund and Johansson, 2024; Cao et al., 2023; Li et al., 2024; Garza et al., 2023), though successful in other fields, are unsuited for producing reliable ecological insights or testing management interventions. For instance, work on multimodal LLMs for environmental prediction (Li et al., 2024) achieves impressive accuracy in forecasting physical variables like streamflow and water temperature, but does not address the mechanistic relationships needed for ecosystem management and applications.

Rather than using AI to replace traditional modelling approaches, recent advances in AI coding capabilities suggest a more promising direction (Xu et al., 2021). Recent demonstrations of LLMs automating scientific processes, from autonomous chemical experimentation (Burger et al., 2023) to biomedical research (Wang et al., 2024), evidence synthesis (Spillias et al., 2024b) and even fully automated scientific discovery (Kramer et al., 2023), highlight their potential for systematic scientific work. The key challenge lies

in developing frameworks that can systematically harness these capabilities while ensuring scientific rigor and maintaining human oversight in the discovery process (Kramer et al., 2023; Spillias et al., 2024a). To the best of our knowledge, such an approach has not been attempted yet in ecological modelling.

To address this challenge, we present “LLM-Enabled Mechanistic Modelling for ecosystem Assessment” (LEMMA), a novel framework that requires minimal inputs, only time-series data and research questions, and aims to produce ecologically sound mathematical models that explain the time-series data (Figure 1). LEMMA addresses the inverse problem of inferring ecologically meaningful mechanistic models and parameters that causally explain observed data. The inferred models can then be used to test management interventions and scenarios. When solving inverse problems through numerical optimization, practitioners usually predetermine the model and its allowable parameter ranges, which defines both the parameter space and its mapping to the observation space. Less commonly, the model itself is inferred directly from observations, as in Scientific Machine Learning (SciML) (Willard et al., 2022), System Identification (Chiuso and Pillonetto, 2019), and Automated Algorithm Discovery (Blazek et al., 2024). Successful model reconstruction in such scenarios requires extensive data to ensure statistical relations are represented in the observations. In addition, a common challenge across all inverse problem approaches is the need to impose constraints that prevent numerically accurate yet empirically unrealistic solutions.

LEMMA differs from traditional and SciML approaches by leveraging the broad scientific and ecological knowledge encoded in large language models (LLMs) through pretraining on diverse sources such as peer-reviewed literature, textbooks, and code repositories. This embedded knowledge is used to impose ecologically meaningful constraints on both model structures and parameter ranges. In addition, LEMMA incorporates a Retrieval-Augmented Generation (RAG) process to identify parameter values and plausible bounds from trusted local repositories and online scientific sources, ensuring biologically realistic parameterization. Within this framework, LEMMA employs Template Model Builder (TMB; an R/C++ library for efficient parameter estimation in complex nonlinear models) as its computational backbone, providing a rigorous statistical foundation for ecological modelling. The system operates iteratively: an LLM proposes candidate model structures as TMB-compatible equations, which are evaluated against time-series data using normalized objective functions (where lower values indicate better fit). These models then undergo evolutionary optimization across multiple generations, with high-performing structures (“individuals”) retained and refined while underperforming models are discarded. This evolutionary process system-

atically yields increasingly accurate representations of ecosystem dynamics and enables exploration of alternative mechanistic hypotheses with quantified support.

We test LEMMA’s ability to construct plausible ecological models through two complementary case studies that test different aspects of ecological modelling, each involving both dependent variables (state variables predicted by the model) and forcing variables (external drivers affecting the system). First, we evaluate the framework’s ability to recover fundamental ecological understanding using synthetic data generated from a well-established nutrient-phytoplankton-zooplankton (NPZ) model (Edwards and Brindley, 1999). This experiment tests LEMMA’s equation-learning capabilities by comparing discovered equations against known mathematical relationships that represent core ecological processes. For our NPZ case study, we additionally evaluate models using ecological accuracy scores (on a scale of 0-8), which measure how well the generated models recover known ecological mechanisms from the reference model. These scores assess specific components like nutrient uptake, phytoplankton growth, and zooplankton dynamics, with higher scores indicating closer alignment with established ecological understanding. Second, we assess LEMMA’s ability to provide management-relevant predictions using synthetic data based on Crown-of-Thorns starfish (COTS) populations on the Great Barrier Reef, derived from existing MICE models (Morello et al., 2014; Rogers and Plagányi, 2022; Plagányi et al., 2014; Condie et al., 2021). This case study included three dependent variables (COTS abundance, fast-growing coral cover, and slow-growing coral cover) and two forcing variables (temperature and larval immigration), testing the framework’s robustness in a complex predator-prey system. We implemented a factorial design comparing three state-of-the-art LLMs, GPT-5, Sonnet-4.5, and Gemini-Pro-2.5, across both case studies, with multiple replicates and evolutionary runs. This comparative approach allowed us to evaluate differences in convergence, predictive accuracy, and ecological realism, demonstrating that LEMMA can both rediscover theoretical relationships and approximate expert-level performance.

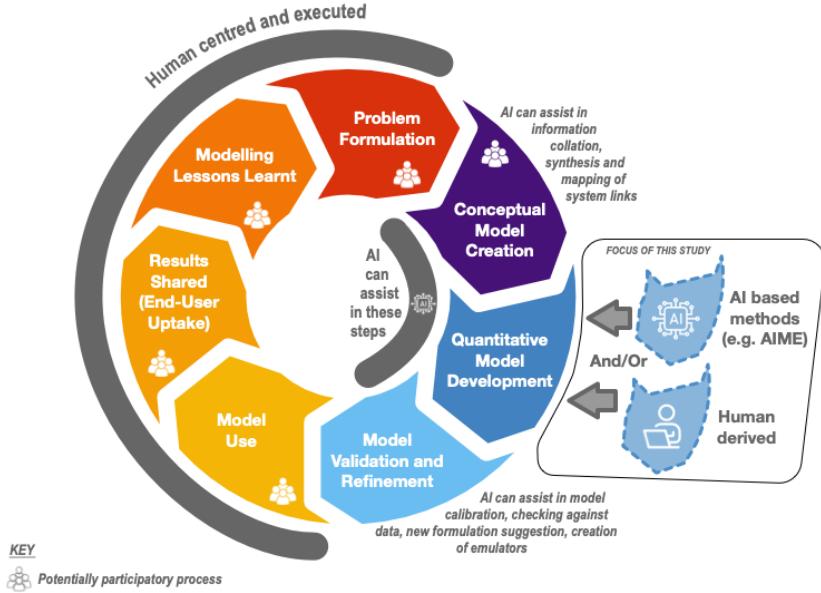


Figure 1: Stylised representation of the iterative modelling process that LEMMA aims to support. Whilst human experts drive the majority of the process, we show that AI-driven processes could play an important role in the Model Development stage.

## 2 Methods

At its core, LEMMA integrates Large Language Models (LLMs) for generating and modifying model structures, Template Model Builder (TMB) for statistical parameter estimation, and evolutionary algorithms for systematic model improvement. All of the code and data underpinning this study are available at the Github repository: <https://anonymous.4open.science/r/EMs-with-LLMs-10E8>.

### 2.1 LEMMA Framework

#### 2.1.1 Model Generation and Improvement

LEMMA uses LLMs to write and modify computer code through Aider (Gauthier, 2024), which is a coding assistant that can create, modify, and interpret local files. Aider can be used in the command-line or called within python scripts, as we have done here, and can receive text and/or images as input, depending on whether the underlying LLM is 'multi-modal' (i.e., can interpret text and images). Each model instance, referred to as an indi-

vidual in our evolutionary framework, consists of three components: (1) a TMB-compatible dynamic model written in C++ that implements a system of equations, (2) a parameters file containing initial values and bounds, and (3) a documentation file explaining the ecological meaning of the model equations (see Section ?? for the complete prompt).

The LLM generates initial parameter estimates for pre-testing model structure before optimization begins. For each parameter, it assigns a priority number that determines optimization order, following established practices in ecosystem modelling (Plagányi et al., 2014).

During non-initial steps, if multi-modal (i.e. can receive images as input) the LLM analyzes performance plots comparing predictions to historical data, otherwise the LLM receives a structured file showing the model fit residuals. After interpreting the model fit, LEMMA makes targeted, ecologically meaningful changes to model equations, implementing one modification at a time to maintain transparency and traceability of successful modelling strategies (see Section ??).

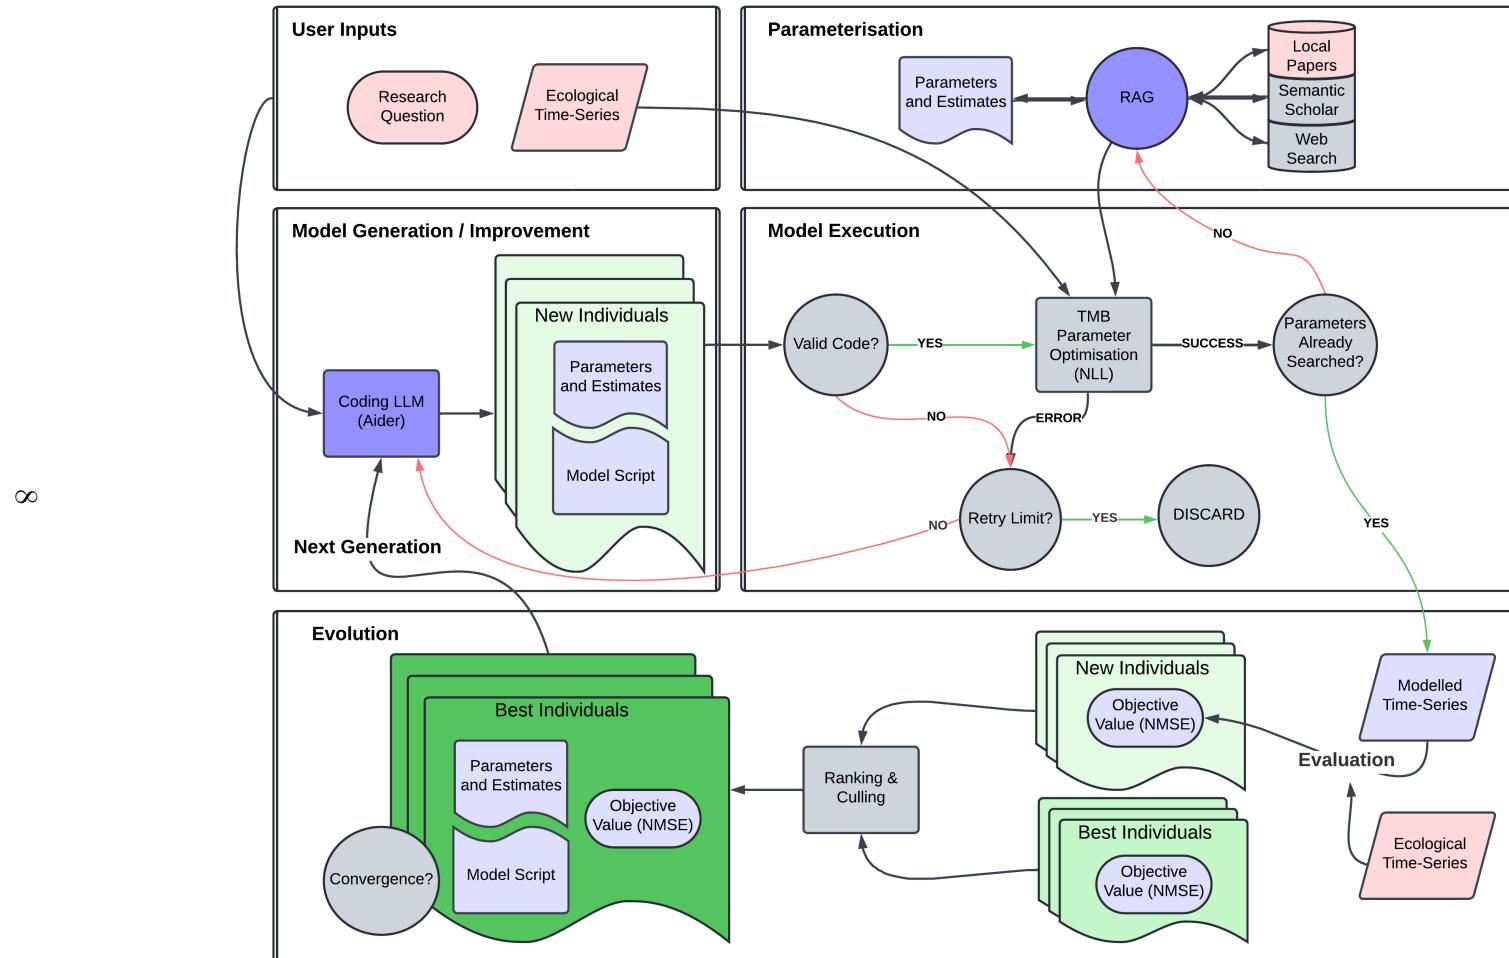


Figure 2: Conceptual diagram of the automated ecological modelling framework, LEMMA. The workflow consists of five main components: (1) User Inputs, where research questions and ecological time-series data are provided; (2) Parameterisation, utilizing RAG-enhanced literature search to estimate parameter values; (3) Model Generation/Improvement, where the Coding LLM creates new individuals with model scripts and parameters; (4) Model Execution, where the LLM’s model code is implemented and TMB is used to optimise parameter values; and (5) Evolution, which evaluates model performance through individual assessment, error handling, and ranking-based selection.

### 2.1.2 User Inputs

The LEMMA framework requires only minimal user input to initiate the modelling process. At the outset, the user provides a short natural language description of the ecological system and research question, which guides the large language model (LLM) in generating ecologically relevant model structures. Two time-series data files are also supplied: a response file containing the dependent variables (state variables) to be predicted by the model, and a forcing file containing external drivers that influence system dynamics. Column headers in these files define the variables that are directly observed and available for model fitting. Importantly, the framework is not restricted to these variables; the LLM may introduce additional latent variables or intermediate processes when such constructs are ecologically justified and improve model performance. This flexibility allows LEMMA to explore a broader range of mechanistic hypotheses than would be possible if limited to observed variables alone.

In addition to these core inputs, users may optionally specify configuration parameters that control the evolutionary process and LLM behaviour, including the sampling temperature, the number of individuals per generation, the number of generations, the convergence threshold, and the proportion of data allocated to training versus testing. Users also select the LLM for code generation, the model for literature retrieval, and the embedding model for semantic search. Finally, the user can specify a directory with curated ecological literature from which LEMMA will identify plausible parameter values and bounds. These inputs collectively define the modelling context while allowing LEMMA to operate with minimal manual intervention and substantial flexibility in model structure.

### 2.1.3 Parameterisation

Upon initialization, the LLM estimates parameter values for each parameter in its proposed model. LEMMA uses this initial estimation to quickly run the model using TMB to identify structural or syntactical errors in the LLM-generated code. If a candidate model is successful in compiling and running, LEMMA refines these estimates using evidence from the scientific literature. Building on the success of LLM-based extraction from ecological literature (Keck et al., 2025; Spillias et al., 2024b), LEMMA implements a RAG architecture to search scientific literature (see Section ?? for detailed RAG implementation). Without this initial estimation step, LEMMA risks wasting time and computational resources searching for parameter values to populate equations that may be structurally or syntactically flawed.

The RAG process works as follows: First, LEMMA prompts an LLM to create detailed semantic descriptions of each parameter, expanding beyond the basic descriptions provided by the coding LLM. For example, if the coding LLM defines a parameter as “growth rate of phytoplankton,” the RAG system might expand this to “maximum specific growth rate of marine phytoplankton in nutrient-rich conditions, measured per day.” These enhanced descriptions aim to improve the relevance of search results when querying literature databases.

To find appropriate parameter values, the RAG system employs a structured, multi-source search strategy. It searches a curated local collection of scientific papers (see Section ??), queries the Semantic Scholar database (Allen Institute for AI, 2024), and performs general web searches through the Serper API (Serper, 2024). LEMMA combines results from all three sources to build a comprehensive understanding of each parameter’s possible values and ecological meaning.

The RAG system then uses LLMs to extract numerical values from the search results, determining not only parameter values but also their valid ranges. The prompt instructs the LLM to identify minimum, maximum, and typical values for each parameter, along with their units and citation information (see Section ??). All parameter information is stored in a structured database that includes bounds, units, and citations.

#### 2.1.4 Model Execution, Optimisation, and Error Handling

**Model Execution.** Models are executed through TMB (Kristensen, 2014), which uses automatic differentiation to efficiently compute gradients for complex, non-linear optimisation problems. The dynamically generated TMB template defines the objective function internally; LEMMA treats this as a black box and does not modify its structure. We use a Control File in *R* which provides the data, starting values, and explicit box constraints for optimisation.

**Optimisation Strategy.** Parameter estimation proceeds in phases: parameters are grouped by LLM-assigned priority, and only those assigned to the current phase are estimated while others remain fixed. After all phases are complete, a final optimisation is performed with all parameters free. This phased approach improves stability and convergence for high-dimensional models. For this proof-of-concept, all parameters are treated as estimable in the optimisation process. The optimiser used is `nlminb`, which minimises the objective function exposed by the TMB template. Lower and upper bounds for each parameter are applied as explicit box constraints at the optimiser

level. Starting values and bounds are constructed from literature-derived values when available; otherwise, LLM-estimated fallbacks are used. Before optimisation, starting values are clamped to the admissible box, and pathological intervals are normalised (swapping if  $\text{lower} > \text{upper}$  and expanding zero-width intervals by a small  $\varepsilon$ ). These checks ensure numerical stability and prevent infeasible parameter configurations.

The optimisation is not iterative in epochs; instead, it uses a gradient-based solver that terminates when one of its convergence criteria is met: (a) gradient norm below tolerance; (b) parameter updates below tolerance; or (c) maximum iterations reached (default in `nlminb` is 150 iterations per phase). These stopping conditions are controlled internally by `nlminb` and are not manually tuned.

Because the TMB template is dynamically generated, LEMMA does not impose a fixed likelihood structure but assumes that the function represents a valid statistical model. The optimisation uses a maximum likelihood framework, leveraging TMB’s automatic differentiation to compute gradients. No regularisation terms are added by LEMMA; the optimisation is unconstrained beyond the explicit parameter bounds. LLM- and literature-provided values are used only to set starting points and bounds, not as penalties in the objective function.

**Error Handling.** Because LLMs often make trivial mistakes in their outputs, we developed an error handling system to address common issues and preserve LLM-progress, rather than discarding partial attempts. For example, on occasion, the LLM coder will attempt to create a system of equations with circular logic (“data leakage”), or will omit a key output variable with a corresponding time-series. To prevent this, we implement code validation checks to ensure that the submitted model is properly formatted and free from logical inconsistencies. For models that fail, LEMMA addresses compilation errors through automated analysis of error messages and implementation of appropriate fixes. For numerical instabilities, LEMMA employs progressive simplification of model structure while maintaining ecological relevance. Each model variant receives up to a user-specified number of fixes (five for our case-studies), with later iterations prompting for simpler model structures that can be iteratively improved. The specific prompts used for error handling are provided in Section ??.

### 2.1.5 Model Evaluation

For each response variable  $j$ , we calculate a normalized mean squared error:

$$\text{NMSE}_j = \begin{cases} \frac{1}{n} \sum_{i=1}^n \left( \frac{y_{ij} - \hat{y}_{ij}}{\sigma_j} \right)^2 & \text{if } \sigma_j \neq 0 \\ \frac{1}{n} \sum_{i=1}^n (y_{ij} - \hat{y}_{ij})^2 & \text{if } \sigma_j = 0 \end{cases} \quad (1)$$

where  $y_{ij}$  represents observed values for variable  $j$  at time  $i$ ,  $\hat{y}_{ij}$  represents corresponding model predictions,  $\sigma_j$  is the unbiased standard deviation of the observed values for variable  $j$  (calculated with  $n - 1$  denominator), and  $n$  is the number of observations. The final objective function value is the mean across all response variables:

$$\text{Objective} = \frac{1}{m} \sum_{j=1}^m \text{NMSE}_j \quad (2)$$

where  $m$  is the number of response variables. NMSE was selected as the primary objective function because it normalizes error by the standard deviation of each time-series, ensuring that all variables contribute equally to model evaluation. This avoids bias toward variables with larger magnitudes and supports balanced optimization across multivariate ecological datasets. For simplicity in this proof-of-concept, we did not weight the time-series in the objective function, however this might prove useful in future work to prioritize uncovering key dynamics.

### 2.1.6 Evolutionary Algorithm Implementation

LEMMA maintains a population of model instances, which we refer to as ‘individuals’, where each individual represents a complete model implementation including its equations, parameters, and performance metrics. Within each generation, individuals undergo parameter optimization using Template Model Builder to find optimal parameter values for their current model structure.

After parameter optimization, individuals are evaluated based on their prediction accuracy. Those achieving the lowest prediction errors (objective values) are selected to become parents for the next generation, while less well-performing individuals are culled and non-functioning ones (those that fail to compile or execute) are discarded.

At the beginning of each new generation, LEMMA creates new individuals in two ways: by making targeted structural modifications to the best-performing parent individuals from the previous generation, and by creating entirely new individuals from scratch when there are not enough functioning individuals. Both of these are done via Aider’s underlying tool handling routines that enable file modifications and creations and are determined by

the underlying LLM’s response to the standard prompts, the project topic, and the current model performance (if modifying an existing model).

## 2.2 Validation Experiments

We conducted two complementary validation case studies of LEMMA. The first validation experiment aimed to see if LEMMA could recover known model equations from synthetic time-series data, whilst the second validation experiment examined real-world applicability through modelling ecological dynamics on the Great Barrier Reef. In addition to the primary experimental design (training on the full dataset), we conducted a supplementary temporal hold-out evaluation for the COTS case study using the best-performing LLM, with three replicate populations, to assess out-of-sample predictive performance.

### 2.2.1 LLM Configuration and Experimental Design

We utilized three leading-edge LLMs to drive LEMMA’s model generation and improvement processes: GPT-5, Sonnet-4.5, and Gemini-Pro-2.5. These models were selected as they represent the current state-of-the-art in language model capabilities at the time of analysis. All models were accessed via OpenRouter, and the analysis was conducted in October 2025. To maintain consistency and reduce complexity in outcomes, we employed GPT 4.1 as a dedicated support model for all parameter estimation and RAG-related queries across all three main LLMs. This approach ensured that differences in model performance could be attributed to the core model generation capabilities rather than variations in parameter retrieval or literature search processes.

To systematically evaluate LEMMA’s performance across different LLMs and ecological systems, we implemented a balanced factorial experimental design. The design consisted of three replicates for each combination of the three LLMs and the two ecological case studies (NPZ and COTS). This yielded a target of 18 populations ( $3 \text{ LLMs} \times 2 \text{ case studies} \times 3 \text{ replicates}$ ).

Each population was initialized with four individuals and allowed to evolve for a target of ten generations, unless convergence was reached earlier (defined as achieving an objective value improvement of less than 0.01 for three consecutive generations).

### 2.2.2 Retrieving Model Equations – NPZ Case Study

We conducted a controlled experiment using synthetic time-series data generated by a well-established nutrient-phytoplankton-zooplankton (NPZ) model from Edwards and Brindley (1999), whose dynamics are well-studied (Boschetti, 2008, 2010). The complete system of equations is presented in Section ?? of the Supplementary Information. This validation tested our framework’s ability to rediscover established ecological relationships from synthetic data where the underlying equations of a system are known, providing a rigorous assessment of the system’s equation-learning capabilities.

In addition to monitoring the convergence of LEMMA’s modelled time-series towards the provided time-series data, we evaluated the framework’s ability to recover nine key ecological characteristics from the original model. These characteristics were organized by equation: three components in the nutrient equation ( $dN/dt$ : uptake, recycling, mixing), four components in the phytoplankton equation ( $dP/dt$ : growth, grazing loss, mortality, mixing), and two components in the zooplankton equation ( $dZ/dt$ : growth, mortality).

During evolution, for each ‘best performer’ in a generation, we used GPT-5-mini to evaluate each model component using a 4-level ordinal scoring system designed to be interpretable and verifiable by ecological experts:

- TRUTH\_MATCH: The mathematical structure is equivalent to the ground truth NPZ model
- ALTERNATE: The implementation matches a recognized alternate formulation from the ecological literature (e.g., different functional response curves from (Franks, 2002))
- SIMILAR\_NOT\_LISTED: The implementation plays the same ecological role but uses a form not represented in the ground truth or literature catalog
- NOT\_PRESENT\_OR\_INCORRECT: The ecological component is missing or cannot be identified

Each characteristic was assigned a weight based on its contribution to its parent equation (e.g., the three nutrient equation components each had weight 0.333, the four phytoplankton components each had weight 0.25, and the two zooplankton components each had weight 0.5). The total ecological score was calculated as the weighted sum of individual scores, then normalized to a 0-1 scale by dividing by the maximum possible score. The complete evaluation prompt with detailed scoring criteria and catalog of alternate formulations is provided in Section ???. This additional evaluation allowed us

to better understand whether objective value improvements were indeed related to improved ecological understanding, or whether they were instead related to spurious mathematical relationships with limited ecological basis. We ran this evolutionary process in four individuals for 10 generations and a convergence threshold of 0.05.

### 2.2.3 COTS Case Study

The Crown-of-Thorns starfish (COTS) case study examined real-world applicability through modelling populations of COTS and their prey, coral, on the Great Barrier Reef. This case study also made two external forcing time-series available to LEMMA, sea-surface temperature and COTS larval immigration quantities. We tested the three main LLMs (GPT-5, Sonnet-4.5, and Gemini-Pro-2.5) and evaluated LEMMA’s ability to match the model created by a human expert in the same context.

The COTS model that we used as a human-derived benchmark for evaluating LEMMA’s outputs was originally developed to specifically evaluate management interventions (Morello et al., 2014). Subsequent model versions and variations have yielded insight into management under environmental perturbations (Rogers and Plagányi, 2022; Condie et al., 2021), derivation of management thresholds (Plagányi et al., 2020; Rogers et al., 2024) and their dynamic implementation (Rogers et al., 2023). Each application differed depending on the objectives and data availability, and how the models were resolved, which required human determinations as to what was included, how it was included, and how it linked with other system aspects where necessary. Simply put, human experts were required to link management objectives and available data to resolve the necessary system aspects for informing specific management actions.

Here we test LEMMA’s ability to develop a model that captures the dynamics of corals and COTS during a COTS outbreak. This not only required LEMMA to link the ecological dynamics but also interpret the life history characteristics of COTS to explain the observed data. Due to time and cost constraints we only performed limited tests using each LLM, where we initialized populations of four individuals for ten generations each. We calibrated these population parameters by balancing the cost of running an individual population against the rate of convergence that we found in our initial tests of the system.

We tracked several key performance metrics for each population:

- Runtime performance: Total runtime and per-generation computation time

- Error resolution: Number of iterations required to achieve successful model implementation in each generation
- Model stability: Proportion of successful, culled (underperforming), and numerically unstable models per generation

We also analyzed the evolutionary trajectories of successful models by tracking their lineage from initial to final states, documenting the frequency and magnitude of improvements across generations. This included measuring the number of generations required to reach best performance and the proportion of attempts that resulted in improved models.

We performed a temporal hold-out evaluation by partitioning the time series into a training period (pre-2000, approximately 70%) used for parameter estimation and a test period (2000-2005; 30%) used once for out-of-sample assessment. For each ecosystem component (COTS abundance, fast-growing coral cover, and slow-growing coral cover), we calculated root mean square error (RMSE), mean absolute error (MAE), and  $R^2$  values to quantify prediction accuracy. By comparing these metrics against those of the human-developed reference model, we could assess whether our automated approach could match expert-level performance in a real-world ecological application.

## 3 Results

### 3.1 Retrieving Model Equations – NPZ Case Study

Analysis of the NPZ validation study revealed that LEMMA achieved strong success in recovering both quantitative fit and ecological structure. The best-performing GPT-5 population reached an objective value of 0.0035 within three generations while maintaining a normalized ecological score of 0.676 (total score = 6.08). One other population converged (objective = 0.0069; normalized ecological score = 0.731), and several additional populations achieved objectives below 0.04 (e.g., 0.0233 and 0.0348), indicating robust predictive accuracy across multiple runs.

Ecological fidelity was consistently high across LLM families. Normalized ecological scores ranged from 0.500 to 0.898, with GPT-5 producing the highest ecological score observed (0.898; total score = 8.08). A score of 0.898 indicates that the model recovered approximately six mechanisms perfectly and provided strong alternatives for the remaining three. Median ecological scores were similar across families (Sonnet-4.5: 0.565; Gemini-2.5-pro: 0.620; GPT-5: 0.731), suggesting that most models captured key ecological processes such as phytoplankton growth and zooplankton dynamics.

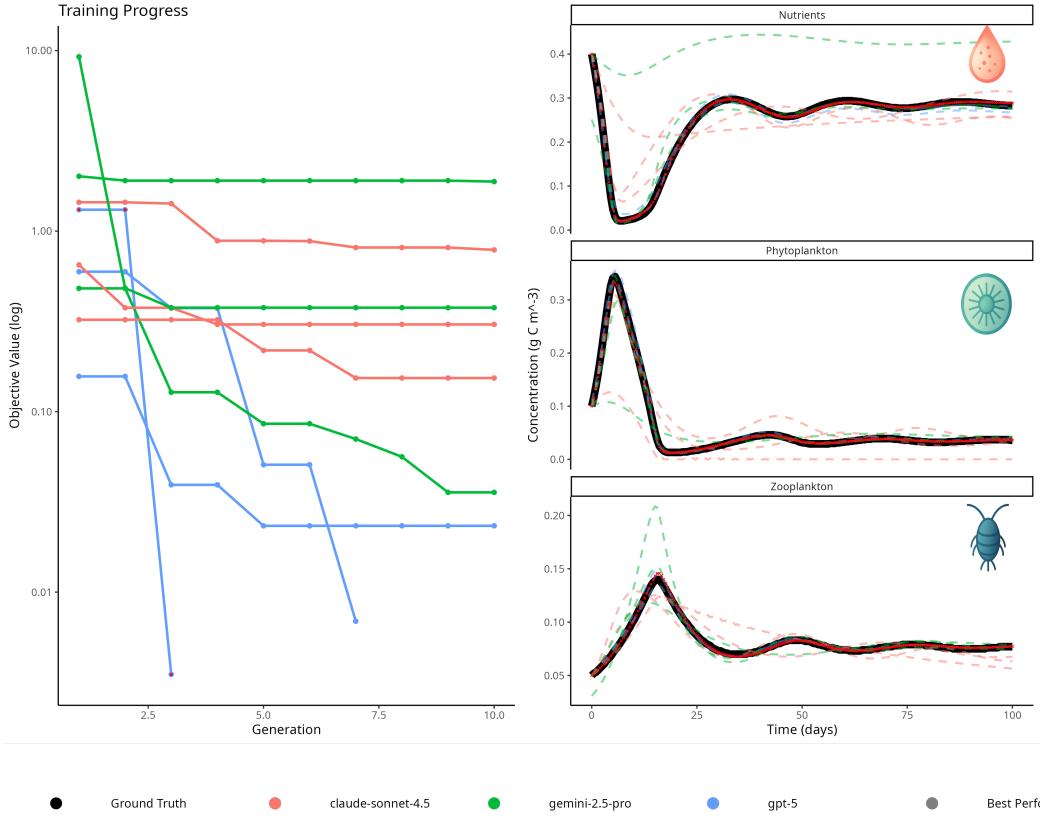


Figure 3: Performance and ecological dynamics of NPZ model retrieval. (A) Training progress showing objective value trajectories across generations (log-scale) for GPT-5, Sonnet-4.5, and Gemini-2.5-pro. GPT-5 achieved the lowest objective value (0.0035) within three generations. (B) Time-series comparison of nutrient, phytoplankton, and zooplankton concentrations ( $\text{g C m}^{-3}$ ) between ground-truth (black lines) and predictions from the best model by objective value (crosses). The model successfully reproduced bloom timing, nutrient drawdown, and trophic phase relationships, indicating strong alignment with ecological mechanisms.

A detailed mechanism-level analysis revealed substantial variation in recovery success across NPZ processes. None of LEMMA’s models perfectly recovered all of the specific ecological mechanisms in the source model from Edwards and Brindley (1999), however all of the mechanisms were recovered at some point by at least one of the individuals (Figure 4). GPT-5 consistently outperformed other LLM families, achieving the highest counts of

truth matches (score = 3) for mechanisms such as phytoplankton growth and zooplankton mortality. Sonnet-4.5 and Gemini-2.5-pro often produced alternative or non-standard formulations (scores = 1-2), particularly for nutrient recycling and phytoplankton grazing loss. Mechanisms related to mixing and recycling were the most challenging across all models, with frequent absent or mistaken formulations (score = 0). These patterns suggest that while overall ecological fidelity was high, certain processes, like phytoplankton and nutrient mixing were often difficult (but not impossible) for LLMs to reconstruct accurately.

The best-performing models achieved objective values as low as 0.0035, demonstrating strong predictive accuracy while maintaining meaningful ecological structure. A detailed analysis of individual ecological characteristics (see Figure ?? in Supplementary Materials) revealed that some mechanisms were more readily recovered than others.

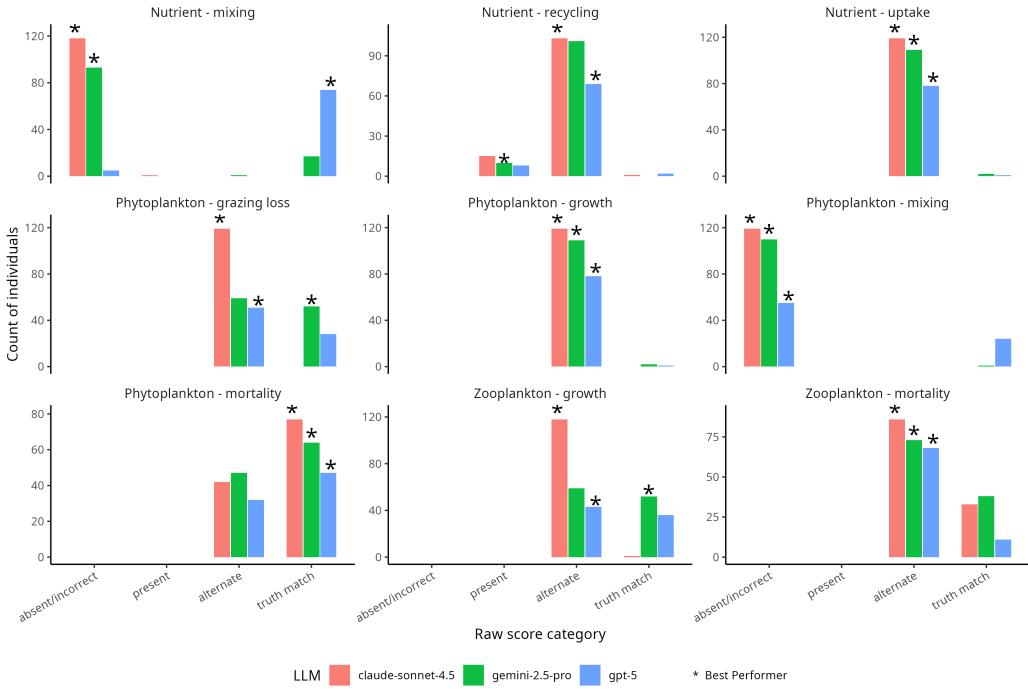


Figure 4: Mechanism-level recovery of NPZ equations across LLM families. Each panel shows the distribution of raw scores for a specific ecological mechanism, where 0 = absent, 1 = present (non-standard), 2 = alternative version, and 3 = truth match. Colors represent different LLM families: Sonnet-4.5, Gemini-2.5-pro, and GPT-5. GPT-5 achieved the highest frequency of truth matches for key processes such as phytoplankton growth and zooplankton mortality, while mechanisms involving nutrient recycling and mixing were most frequently absent or represented by alternative formulations.

A process-level comparison of fluxes revealed that LEMMA successfully reproduced most ecological dynamics of the NPZ system (Figure 5). Uptake, phytoplankton growth, and zooplankton mortality exhibited near-perfect alignment with the ground-truth model, indicating that core trophic interactions were accurately reconstructed. Grazing loss and recycling processes were also well captured, though minor deviations were observed in recycling to nitrogen at higher flux values. In contrast, mixing-related fluxes (nutrient and phytoplankton mixing) were absent in the LEMMA-derived model, consistent with the mechanism-level analysis showing these processes as the most challenging to recover. Despite these gaps, the overall fidelity of flux dynamics underscores LEMMA’s ability to infer mechanistic structure beyond

simple curve fitting.

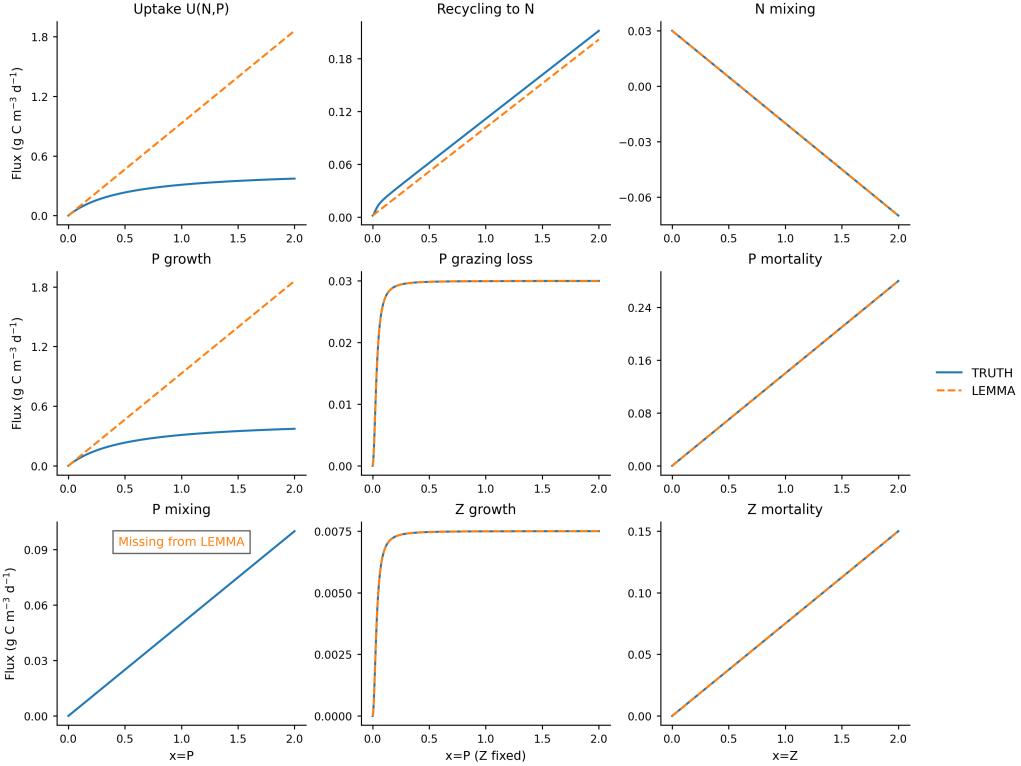


Figure 5: Comparison of process-level fluxes between the ground-truth NPZ model and the best-performing LEMMA-derived model (GPT-5). Each panel shows fluxes ( $\text{g C m}^{-2} \text{ d}^{-1}$ ) for key ecological processes: (1) Uptake ( $U(\text{IN},\text{P})$ ), (2) Recycling to N, (3) N mixing, (4) Phytoplankton growth, (5) Phytoplankton grazing loss, (6) Phytoplankton mortality, (7) Phytoplankton mixing, (8) Zooplankton growth, and (9) Zooplankton mortality. Solid blue lines represent the ground-truth model, while dashed orange lines represent the LEMMA model. Strong alignment is observed for most processes, particularly uptake, phytoplankton growth, and zooplankton dynamics, while discrepancies appear in mixing-related fluxes where LEMMA did not recover explicit formulations.

### 3.2 COTS Case Study

For all tested LLMs, LEMMA was able to generate ecosystem models with prediction accuracy approaching the quantitative fit of the expert-developed

model (objective value NMSE: 0.2312). After ten generations, objective values across LLMs (based on three independent runs per model) were as follows: GPT-5 achieved the best performance, with objective values ranging from 0.31-0.45 (lower is better), followed by Sonnet-4.5 (0.33-0.44) and Gemini-2.5-pro (0.47-0.69). Component-specific analysis revealed varying levels of prediction accuracy, with models showing strongest performance in predicting fast-growing coral cover and slow-growing coral cover, while maintaining reasonable accuracy for the more volatile COTS abundance patterns (Figure 6). GPT-5 was able to correctly model the outbreaks of COTS, as requested in the initial prompt, while Gemini-2.5 Pro’s best model was out of phase and Sonnet-4.5’s identified the first, but not the second.

The best-performing models generated by GPT-5, Claude Sonnet 4.5, and Gemini 2.5 Pro exhibited differences in structural design and ecological mechanisms, whilst also sharing several elements. GPT-5 aligned most closely with the demographic logic of the human reference model, while Gemini and Claude adopted more simplified, outbreak-oriented frameworks. Across all LLMs, certain core features remained consistent: density-dependent population regulation in crown-of-thorns starfish (COTS), resource limitation driven by coral availability, and differential effects on fast- versus slow-growing coral species. These commonalities occurred alongside substantial variation in life-history representation, recruitment formulation, immigration handling, mortality drivers, predation dynamics, and thermal stress responses (see Supplement ??).

Among the LLM-generated models, GPT-5 most closely resembled the human reference model in population structure. GPT-5 modelled crown-of-thorns starfish (COTS) using two stages, juveniles and adults, while the human model applied a more detailed age-structured approach with three classes and age-specific mortality rates. In contrast, Gemini simplified the population to a single undifferentiated compartment, and Claude used a single adult compartment supplemented by episodic recruitment pulses. For recruitment, GPT-5 incorporated a Beverton-Holt-like taper on adult abundance, modified by environmental factors and lagged immigration, which paralleled the continuous Beverton-Holt density dependence used in the human model. Gemini, on the other hand, linked recruitment to coral consumption and introduced Allee effects and temperature windows without an explicit stock-recruitment curve. Claude adopted a threshold-based pulse mechanism rather than a continuous function. GPT-5 handled immigration by integrating with recruitment and applied a temporal lag, echoing the demographic embedding seen in the human model, which introduced immigration during age-0 formation with lognormal variability. Gemini instead added immigrants directly to the adult compartment, while Claude com-

bined immigration with favorability indices and growth multipliers, creating a compound influence on population dynamics. GPT-5 and Gemini both applied baseline mortality augmented by density-dependent feedbacks, which partially reflected the human model’s approach of age-specific mortality tied explicitly to prey availability. Claude introduced a starvation multiplier instead that amplified mortality when coral cover was low. For predation, GPT-5 blended Holling Type II and Type III responses, producing low-prey refuges, whereas the human model used a sigmoid saturation curve with explicit prey-switching rules. Gemini employed a standard Holling Type II functional response, and Claude incorporated prey preference weighting rather than explicit switching. For coral growth and thermal stress, GPT-5 penalized coral growth exponentially above a thermal threshold and included a heat-loss term, which partially mirrored the human model’s use of Gaussian performance curves for growth combined with logistic bleaching mortality. Gemini restricted thermal effects to logistic bleaching only, while Claude applied linear mortality beyond stress thresholds without Gaussian modulation. Finally, temperature effects on COTS recruitment were absent from the human model but present in all LLM implementations. GPT-5 and Gemini used Gaussian temperature functions to modulate recruitment, whereas Claude embedded temperature influences within outbreak pulses and adult growth factors.

### 3.3 Evolutionary Behaviour

There were notable differences in the behaviour and speed of convergence across all LLMs tested (Figure ??). The Gemini-2.5 Pro models were the fastest overall, completing ten generations in approximately 36-50 minutes, with a mean time per generation of just over four minutes. GPT-5 models showed intermediate performance, averaging about 12 minutes per generation.

Patterns of improvement also differed markedly. Sonnet-4.5 and Gemini-2.5 Pro generally produced well-functioning models early but struggled to achieve consistent gains across generations, with mean improvement rates close to zero and frequent plateaus. GPT-5 demonstrated more stable progress, including the best-performing population overall (objective value 0.0035), though most GPT-5 runs also exhibited diminishing returns after initial improvements.

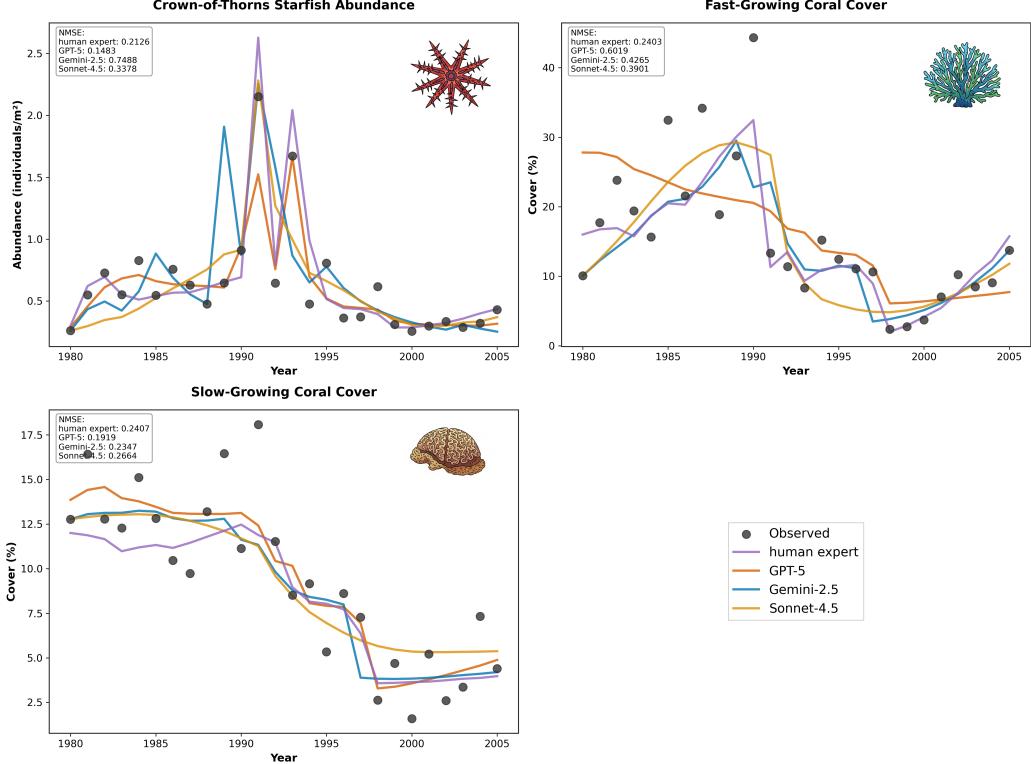


Figure 6: Comparison of model predictions across ecosystem components. The plots show observed versus predicted values for COTS abundance, fast-growing coral cover, and slow-growing coral cover, demonstrating the models’ ability to capture key ecological patterns and relationships. Objective values (obj) shown in the legend represent the normalised mean squared error for all three variables, where lower values indicate better model performance.

### 3.3.1 Temporal Hold-Out

Because GPT-5 demonstrated the strongest overall performance, we used it as the base LLM to test whether the LEMMA framework could create models capable of predicting out-of-sample datapoints. After ten generations, we found that the best-performing model (objective value: 0.29) had a reasonable degree of out-of-sample prediction performance on the withheld test dataset, with particularly strong predictive power for fast-growing coral cover ( $R^2 = 0.81$ , RMSE = 2.22, MAE = 1.92). Interestingly, this was a better objective value than the previous experiments where LEMMA had access to the entire time-series.

For slow-growing coral cover, the model achieved moderate predictive

accuracy ( $R^2 = 0.17$ , RMSE = 1.6, MAE = 1.3), effectively capturing the general declining trend while showing some deviation in precise values. COTS population predictions demonstrated strong accuracy ( $R^2 = 0.75$ , RMSE = 0.07, MAE = 0.05), successfully capturing the timing and magnitude of major outbreaks.

Figure 7 illustrates these prediction capabilities, showing both training period performance (pre-1997) and out-of-sample predictions (1997-2005). The model's ability to maintain consistent error metrics (RMSE and MAE) while capturing both rapid population dynamics and slower coral cover changes suggests it successfully identified fundamental ecological relationships governing this system.

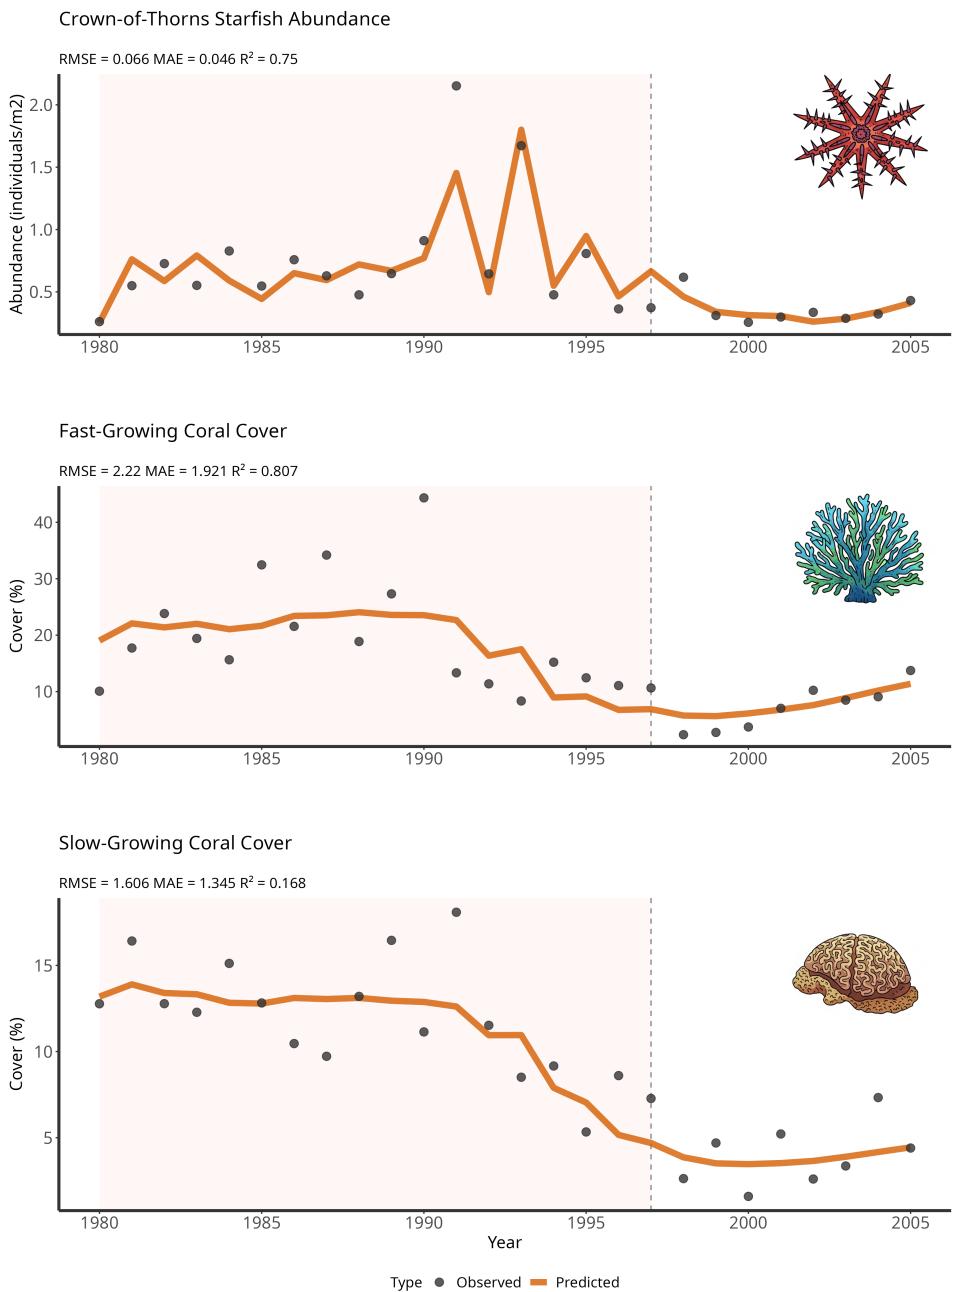


Figure 7: Temporal hold-out evaluation of LEMMA driven by the best-performing LLM (GPT-5) showing predictions against observed data. LEMMA was provided with 70% of the time-series data (pink shaded region) and developed a model that was then evaluated on the full time-series, including the remaining unseen 30% <sup>25</sup> of the time-series data (white region). Top: COTS abundance predictions showing strong capture of outbreak dynamics. Middle: Fast-growing coral cover predictions demonstrating tracking of recovery. Bottom: Slow-growing coral cover predictions illustrating general declining trends. Orange lines represent model predictions, black dots show observed data.

### 3.4 Parameter Retrieval

We assessed how frequently LEMMA linked model parameters to external evidence sources, including Semantic Scholar searches and a locally curated document store. Across all generated models (21 populations, 2,619 parameters), citation integration was modest: only 3.1% of parameters were linked to explicit citations, 1.2% to Semantic Scholar matches, and 2.0% to document store references. The majority of parameters lacked any external linkage. Differences emerged between case studies. NPZ-based populations showed slightly higher citation rates (3.7%) and semantic matches (2.7%) than COTS-focused populations (2.8% and 0.59%, respectively). Document store usage was notably higher for COTS models (2.4%) compared to NPZ (1.1%). This disparity likely reflects the fact that the document store was curated specifically for the COTS case study, whereas no curation was performed for NPZ, limiting retrieval opportunities. At the individual model level, citation coverage rarely exceeded 10% of parameters, and many individuals contained no cited parameters. Where citations were present, they typically supported temperature optima, growth coefficients, and bleaching thresholds, parameters with well-documented ecological relevance. These results indicate that while LEMMA can incorporate literature-based evidence, this capability was underutilized in practice. Enhancing citation integration may require prioritizing literature-grounded parameter selection and expanding curated resources for all ecosystem contexts.

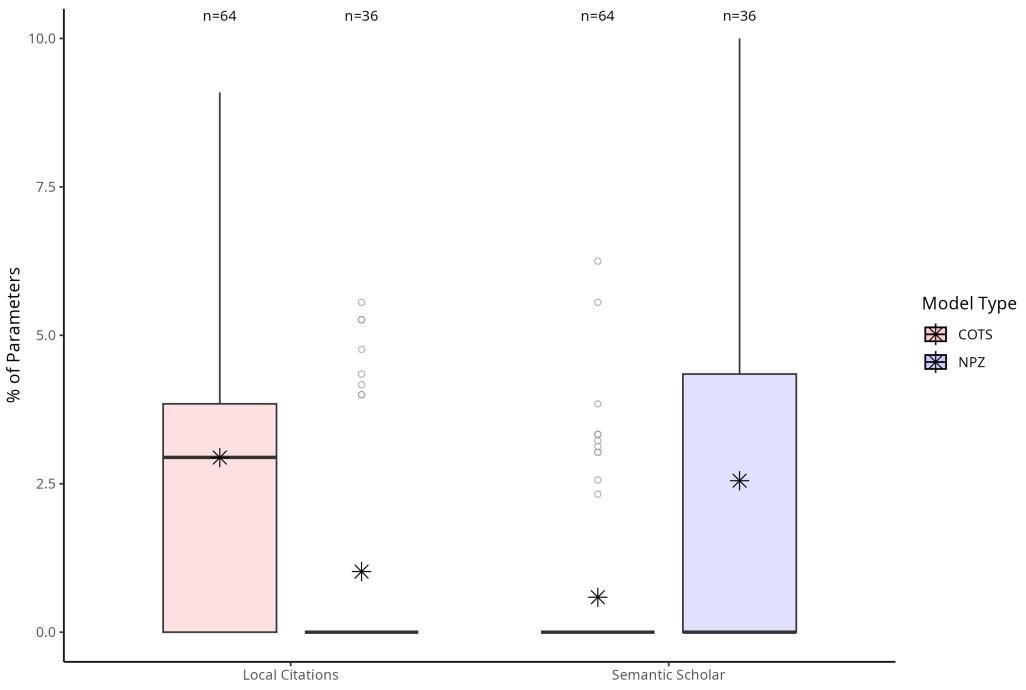


Figure 8: Comparison of citation integration between NPZ and COTS case studies. Boxplots illustrate the proportion of model parameters with citations (Semantic Scholar or local document store). NPZ models were more reliant on Semantic Scholar because the local document store was curated specifically for the COTS case study. Asterisks represent the best performer (lowest nRMSE objective value) for each model type.

## 4 Discussion

Our complementary validation studies demonstrate the viability of AI-driven automation in ecological model development and calibration. The NPZ validation revealed LEMMA’s ability to recover known ecological relationships from synthetic data, with the best-fitting individual achieving a normalized ecological score of 0.676, which is moderately good and reflects recovery of all mechanisms except phytoplankton mixing; the remaining components were either identical to the NPZ reference model or accepted alternatives from Franks (2002) (objective value = 0.0035), while the highest ecological score observed was 0.898 for a different individual with a higher objective error (objective value = 0.650) While LEMMA did not perfectly reconstruct the original equations after ten generations, it successfully identified key mech-

anisms like Michaelis-Menten kinetics and predator-prey interactions. The negative correlation between ecological accuracy and objective values suggests that improvements in model fit were achieved through discovery of correct ecological relationships rather than overfitting. This ability to balance model fit with ecological realism represents an important advance in ecological modelling.

The CoTS case study further demonstrated LEMMA’s practical utility, with LEMMA successfully generating models that approached the predictive performance of human expert models, albeit with less consistency across all variables. While LEMMA’ models captured key ecological dynamics, they did so using structurally simpler formulations, and their performance varied depending on the specific ecosystem component being modeled.

Notably, LEMMA-generated models achieved predictive performance that was broadly similar to that of the human expert model, despite substantial structural differences. These results suggest that, under certain conditions, AI-driven approaches may approximate expert-level outcomes, though further validation and refinement are needed to ensure robustness and ecological fidelity (see Supplementary Materials for detailed comparison). While the human model implemented an age-structured COTS population with explicit age classes and a Beverton-Holt stock-recruitment relationship, LEMMA models generally employed simpler, unstructured population approaches. Similarly, the human model featured an explicit prey-switching function for COTS predation preference between coral types, whereas AI models used various functional responses ranging from coral-dependent reproduction to logistic growth with food limitation. These structural differences highlight an important trade-off: the human model exhibited greater mechanistic detail reflecting domain expertise and input from domain experts, while AI models achieved similar performance with more parsimonious formulations. This is consistent with the coexistence of multiple numerically valid representations of the same ecological system where each offers different insights into the underlying mechanisms (Patterson et al., 2001).

Importantly, our findings also address the “forecast trap” described by (Boettiger, 2022), wherein selecting models primarily on predictive skill can lock decision-makers into policies that are accurate in forecast space but suboptimal in utility space. In capacity-constrained settings, having even a single model is often a luxury, and assembling an ensemble is typically out of reach. We hope that LEMMA, or tools like it, will enable diverse sets of models to be built, compared, and ensembled with relative ease. By evolving interpretable mechanistic candidates, LEMMA broadens the model set and supports robust decision workflows, reducing the risk of converging on a well-calibrated yet policy-misaligned model.

This flexibility also opens the door to hybrid approaches that combine the strengths of both expert-driven and AI-generated models. By explicitly prompting LEMMA to include particular ecological processes, such as age structure, prey-switching behavior, or nutrient mixing, researchers can guide model structure to reflect known system dynamics or stakeholder priorities. This capability enhances the utility of AI-generated models for hypothesis testing, scenario exploration, and applied decision-making, particularly in contexts where certain mechanisms are known to be ecologically or socially important.

An important consideration in automated model generation is the provenance of parameter values. In this study, we restricted parameter retrieval to curated local literature and Semantic Scholar searches to ensure reliance on peer-reviewed information, noting that LEMMA incorporates the ability to also perform general web sources should the user choose. While restricting the source of parameter values improves transparency and scientific rigor, it also means that when literature values were unavailable, the system defaulted to LLM-derived initial estimates, which - though often plausible - lack explicit citation. Future enhancements such as user-defined whitelists and confidence scoring to further strengthen parameter provenance may help to address these limitations.

Unlike traditional machine-learning approaches which use regularisation and early-stopping to avoid overfitting, LEMMA employs two nested optimisation processes, both of which may need mechanisms to address overfitting. For parameter estimation within TMB, literature-informed bounds act as weak regularisation, but no additional penalties are applied. For structural evolution, we limit changes to one modification per generation and cap the number of generations. In the NPZ case study, improvements in fit were negatively correlated with ecological accuracy scores, suggesting that gains were achieved through mechanistically sound changes rather than overfitting. Nonetheless, explicit complexity penalties or information-theoretic criteria (e.g., AIC/BIC) could further reduce overfitting risk in future iterations.

## 4.1 Contrasting Approaches to AI in Ecological Modelling

Recent advances in AI have demonstrated remarkable capabilities in ecological time-series prediction. Studies using transformer architectures and diffusion models, including multimodal approaches like LITE (Li et al., 2024), have shown high accuracy in direct forecasting of environmental variables (Morales-García et al., 2024; Gandhi et al., 2024). While these methods

effectively handle challenges like missing data and distribution shifts, they treat the system as a black box, learning patterns directly from time-series data without explicitly modelling underlying mechanisms. While our study did not directly compare our approach with black-box methods, our NPZ validation study suggests a key advantage of our approach: the ability to provide insights into fundamental ecological processes like nutrient cycling or predator-prey dynamics through explicit model discovery. This interpretability is a theoretical advantage over black-box approaches, though comparative studies would be needed to fully evaluate the relative strengths of each approach in specific ecological contexts.

Our evolutionary approach fundamentally differs by using AI to generate actual ecological models rather than make direct predictions. Instead of training neural networks to forecast future values, LEMMA evolves interpretable models with meaningful parameters that capture real biological and physical processes. This distinction is crucial for several reasons. First, our generated models provide scientific insight into system behavior, revealing mechanisms and relationships that direct prediction approaches typically cannot without careful oversight (Adams et al., 2017). Second, the models maintain biological plausibility through explicit parameter constraints and mechanistic formulations, ensuring their utility for management applications. Third, because they capture fundamental processes rather than just patterns, these models can potentially be transferred to new scenarios and used to explore management interventions. The relationship between our approach and direct prediction methods is nuanced. Recent time-series prediction approaches using transformer architectures have achieved impressive accuracy, with mean squared errors as low as 0.001-0.04 for normalized predictions (Morales-García et al., 2024) and root mean squared errors reduced by up to 52% compared to traditional methods (Gandhi et al., 2024). While our evolved models may not always match these pure prediction accuracies, they offer advantages in interpretability, scientific insight and potential for strategic and tactical applications.

Importantly, these approaches need not be viewed as mutually exclusive. Comparing mechanistic models with black-box predictions can be particularly insightful, especially when the two approaches diverge. For instance, under novel conditions like future climate scenarios, differences in predictions could highlight processes that are not well-captured by mechanistic models or reveal patterns that black-box approaches detect but cannot explain. When predictive approaches outperform mechanistic models, this divergence can guide researchers toward missing processes or relationships that should be incorporated into mechanistic understanding. Our framework demonstrates that it's possible to achieve both reasonable predictive accuracy and mean-

ingful ecological interpretability, with each approach offering complementary strengths.

This focus on model generation rather than direct prediction aligns with the needs of ecosystem-based management, where mechanistic understanding is as important as predictive accuracy. The interpretability of AI-evolved models enables users to assess the credibility of predictions and understand the mechanisms driving system behavior thereby allowing for informed management interventions, advantages not readily available with black-box prediction approaches. Work examining automated scientific discovery emphasizes the importance of maintaining human oversight while leveraging AI’s computational capabilities (Kramer et al., 2023; Spillias et al., 2024a). Our approach directly addresses this need by producing interpretable models that facilitate meaningful human oversight while leveraging AI’s capabilities for systematic exploration of model space. Crucially, this process is not solely computational: models are refined through expert collaboration, which synthesizes domain knowledge and builds confidence and legitimacy in the resulting models. This step is essential for filtering out plausible but ultimately incorrect mechanisms that may arise from AI-driven parsimony. For example, some temperature-related functions affecting COTS in LEMMA-derived models lack empirical support and thus their inclusion would likely be removed during expert-led refinement. While such mechanisms may be ecologically plausible, their inclusion without validation risks misleading management decisions, especially under extreme conditions. Thus, expert review acts as an important filter, ensuring that models are not only interpretable, but also credible and useful for decision-making.

## 4.2 Limitations and Future Directions

Despite promising results, several limitations warrant consideration. The observed variation in convergence rates across populations suggests that initial conditions significantly influence model evolution trajectories. While the best performing population achieved rapid convergence within five generations, other populations required more than twice as many generations to approach similar performance levels. In our current pipeline, convergence is also affected by the phased optimisation schedule and box-constraint handling: parameters are estimated in priority-ordered phases with starts clamped into literature- or LLM-derived bounds and pathological intervals normalised. Although these safeguards improve numerical stability, they can also interact with initial conditions in nontrivial ways (e.g., temporarily parking parameters on bound faces early in the schedule and delaying subsequent improvements).

An important limitation in our current implementation is the treatment of all model parameters as estimable quantities in the optimization process, even when well-established values exist in the literature. While our RAG system successfully retrieves literature-based values and ranges, these are used as initial estimates and bounds rather than fixed quantities. This design choice reflects the reality that literature values may not always be appropriate for the specific ecological context being modeled. Estimating parameters within literature-derived bounds allows the model to adapt to system-specific deviations, offering greater flexibility and ecological realism. Future versions of the framework could distinguish between parameters that truly require estimation and those that could be fixed, reducing the parameter space and improving computational efficiency. This would not only reduce the parameter space for optimization but also better incorporate established ecological knowledge into the modelling process and make the process less resource-intensive when doing calculations.

One limitation identified in this study is the low level of citation integration. Only about 3% of parameters are linked to explicit sources. It is reasonable to expect that not all parameters would have readily available citations because many are scaling constants or other mathematical utility parameters that typically do not appear in the literature. This gap is also unsurprising given that the parameter search in this study relied only on abstracts from Semantic Scholar, while full-text scientific articles are often locked behind paywalls. This situation highlights the importance of human oversight. Experts can build on LEMMA’s outputs by supplementing parameterization with their own efforts, ensuring that critical biological parameters are grounded in empirical evidence while computational parameters remain transparent and well documented. LLMs have been demonstrated to perform well at extracting ecological data from curated databases (Gougherty and Clipp, 2024; Keck et al., 2025), therefore we expect that an important first step before deploying LEMMA would be to provide it with as much high-quality information as possible. Future iterations of LEMMA could address these limitations by incorporating more flexible agentic frameworks that can dynamically access curated databases and apply advanced reasoning to identify high-quality sources during web searches. Such improvements would enable the system to prioritize peer-reviewed literature, apply confidence scoring, and cross-validate retrieved values against multiple sources. These enhancements would strengthen parameter provenance and reduce reliance on uncited estimates generated by language models.

It is important to note that our prompt provided only high-level guidance on ecological realism, numerical stability, and reporting requirements, without specifying how parameter estimation should be implemented in TMB.

Despite this minimal guidance, the LLM consistently generated models that adopt a forward-simulation approach: initializing at the first observation, simulating trajectories through time, and minimizing a trajectory-wide error metric. This convergence on a single strategy suggests that, under generic ecological modelling prompts, LLMs may default to conceptually straightforward but computationally fragile approaches. Future work could explore this outcome by explicitly prompting for alternative estimation paradigms, such as state-space formulations (Auger-Méthé et al., 2021), gradient matching (Ellner et al., 2002), or one-step-ahead prediction objectives (Munch et al., 2023), thereby enabling systematic comparison of inference strategies within the same automated workflow.

There are numerous future avenues for validating, improving, and extending this framework. First, there are several hyper-parameters that likely control the success and speed of convergence of the framework (LLM-choice, LLM temperature setting, number of individuals per generation, prompt construction, etc.). Systematic testing across these choices may reveal optimal configurations for convergence. In particular, the comparative analysis of different AI configurations (as detailed in Section 3.2 and Figure ??) reveals trade-offs between model choice and rate of improvement. While GPT-5 demonstrated the most stable improvement and achieved the best overall performance, Sonnet-4.5 and Gemini-Pro-2.5 often produced strong initial models but did not consistently improve across generations. Future work could explore hybrid approaches that leverage the strengths of different LLMs at various stages of model development, or employ different models consecutively over multiple generations. Further, ongoing testing of new LLMs as they are released may yield considerable gains in efficiency and cost-saving. Second, we have tested a relatively simple ecosystem model with three dependent variable time-series and two forcing variable time-series. Simple systems like these will be limited in real-world utility, and therefore testing on more complex systems with tens or hundreds of time-series will be needed. Incorporating spatial components may also be possible and will greatly improve the utility of this framework. Third, accessing relevant scientific information for the parameter RAG search is limited by the user’s ability to either curate a local database of relevant materials, or access scientific papers online. Fourth, we have demonstrated that it is possible for this LLM-based system to generate multiple, distinct models for a given system. Choosing between similarly performing, but ecologically distinct models may be necessary for experts with ecological knowledge, or perhaps employing approaches that ensemble multiple plausible models may allow for the reduction in uncertainty (Baker et al., 2017; Gårdmark et al., 2013; Vollert et al., 2024). Finally, our evaluation approach relies mainly on error-based metrics, which may not

fully reflect ecological plausibility. Recent work on LLM-based evaluation policy extraction (Cheng et al., 2025) offers a promising direction for incorporating expert-informed, interpretable criteria into automated workflows. Future versions of our framework could integrate such methods to improve ecological relevance in model selection.

### 4.3 Implications for Ecosystem-Based Fisheries Management

The successful application of LEMMA to COTS populations on the Great Barrier Reef demonstrates its potential for use in developing plausible models that can provide insights on pressing ecosystem problems relevant to ecosystem based management. While more development and experience is needed to ascertain AI-derived models directly into decision making contexts, there is certainly sufficient evidence that they can be used as a means of hypothesis generation and for the development of plausible models. This may help to accelerate the model-building process, bringing the process more into line with timeframes pertinent to pressing tactical management questions. The framework's capacity to capture both short-term outbreak dynamics and longer-term ecosystem changes provides scientists supporting managers with valuable insights for intervention planning. The comparable performance between AI-generated models and human expert approaches suggests that automated modelling could complement traditional methods, accelerating the development and evaluation of management strategies.

While the current implementation of LEMMA demonstrates faster turnaround than typical manual workflows, reducing model development from weeks or months to hours, these gains reflect relatively simple case studies and yielded models of lesser complexity than human-generated models. As LEMMA matures and is tested in more complex systems, this enhanced efficiency may create capacity for more timely intervention in response to emerging ecological threats, such as in biosecurity emergencies, where management actions depend on rapid model development and deployment. LEMMA's ability to integrate multiple data sources, both locally and from web search, and account for both biological and environmental factors provides a robust foundation for developing early warning systems and evaluating potential management interventions.

To ensure safe and effective integration of LEMMA into ecosystem-based fisheries management, we recommend a set of best practices and propose that LEMMA would work best embedded within a human-driven workflow (Figure 9).

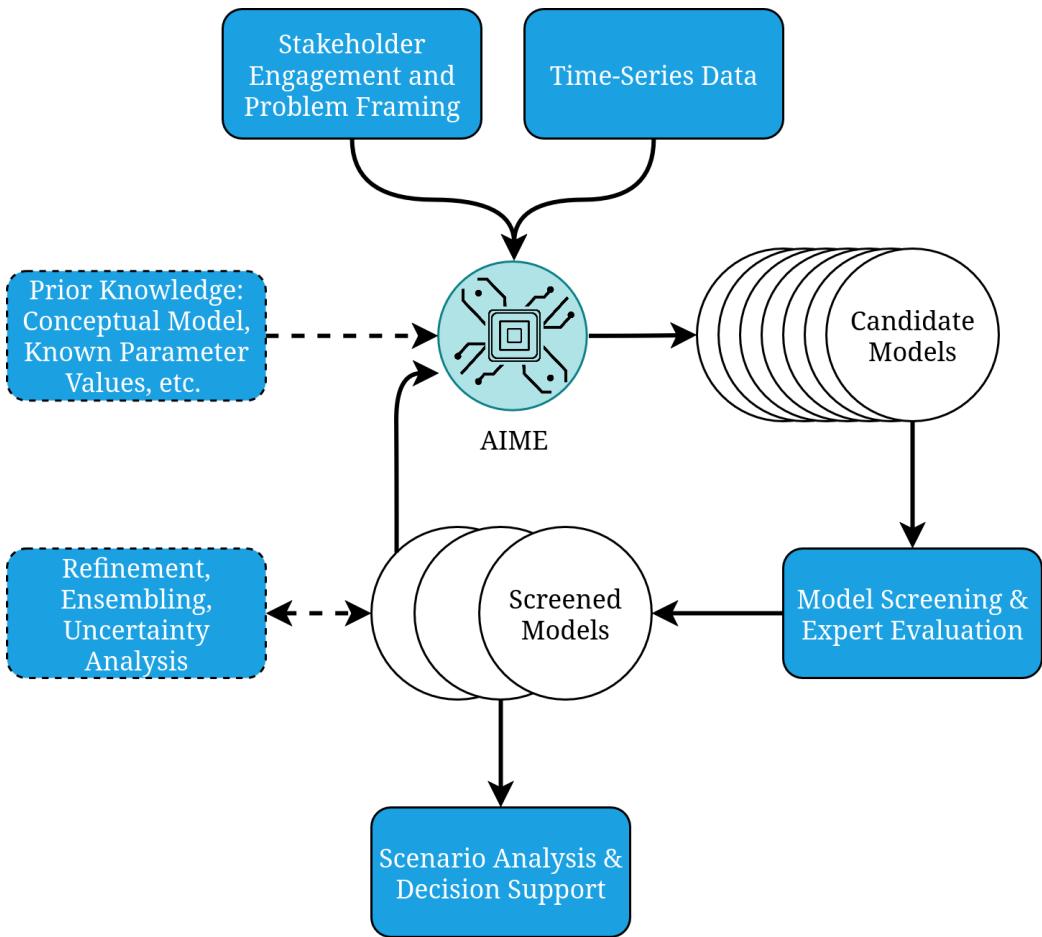


Figure 9: The LEMMA framework workflow integrating human expertise with AI-driven model development. The diagram illustrates how stakeholder engagement, time-series data, and prior ecological knowledge inform the LEMMA process, which generates candidate models that can be evaluated and refined by human experts, ultimately supporting ecosystem management decisions.

1. **Stakeholder Engagement** Begin with stakeholder engagement to inform problem framing and define the ecological questions LEMMA will address. Early involvement ensures models are relevant to management needs and incorporate local knowledge.
2. **Expert Review** All AI-generated models should be reviewed by domain experts to validate ecological plausibility and ensure alignment with management objectives.
3. **Complementary Use** Use LEMMA as a complementary tool that supports - rather than replaces - traditional modelling workflows, especially for rapid prototyping or exploring alternative hypotheses.
4. **Model Transparency** Maintain transparency through clear documentation of equations, parameters, and assumptions to ensure traceability and reproducibility.
5. **Parameter Assessment** Critically assess parameter values sourced from literature. Fix well-established values where appropriate to reduce uncertainty.
6. **Rigorous Validation** Conduct thorough validation, including cross-validation and out-of-sample testing, before applying models to decision-making.
7. **Training and Capacity Building** Provide training to ensure managers and researchers can interpret and apply AI-generated models responsibly.

In conclusion, LEMMA represents an advancement in ecological modelling that bridges the gap between computational efficiency and ecological insight. By dramatically accelerating model development while maintaining scientific rigour, this framework offers a powerful new tool for researchers and managers facing urgent ecological challenges. As environmental pressures intensify globally, the capacity to rapidly develop, test, and deploy ecologically sound models will become increasingly valuable for effective conservation and management of marine ecosystems.

## References

Adams, M. P., Collier, C. J., Uthicke, S., Ow, Y. X., Langlois, L., and O'Brien, K. R. (2017). Model fit versus biological relevance: Evaluat-

- ing photosynthesis-temperature models for three tropical seagrass species. *Scientific reports*, 7(1):39930.
- Allen Institute for AI (2024). Semantic scholar academic graph api.
- Auger-Méthé, M., Newman, K., Cole, D., Empacher, F., Gryba, R., King, A. A., Leos-Barajas, V., Mills Flemming, J., Nielsen, A., Petris, G., et al. (2021). A guide to state-space modeling of ecological time series. *Ecological Monographs*, 91(4):e01470.
- Baker, C. M., Gordon, A., and Bode, M. (2017). Ensemble ecosystem modeling for predicting ecosystem response to predator reintroduction. *Conservation Biology*, 31(2):376–384.
- Blazek, P. J., Venkatesh, K., and Lin, M. M. (2024). Automated discovery of algorithms from data. *Nature Computational Science*, 4(2):110–118.
- Boettiger, C. (2022). The forecast trap. *Ecology Letters*, 25(7):1655–1664.
- Boschetti, F. (2008). Mapping the complexity of ecological models. *Ecological Complexity*, 5(1):37–47.
- Boschetti, F. (2010). Detecting behaviours in ecological models. *Ecological Complexity*, 7(1):76–85.
- Burger, B., Kunkel, P., Koopman, R., Bickelhaupt, F. M., Schlüter, F., and Schwab, P. (2023). Autonomous chemical research with large language models. *Nature*, 624:505–512.
- Bylund, D. and Johansson, V. (2024). ChatGPT for Time Series Predictions.
- Cao, D., Jia, F., Arik, S. O., Pfister, T., Zheng, Y., Ye, W., and Liu, Y. (2023). Tempo: Prompt-based generative pre-trained transformer for time series forecasting. *arXiv preprint arXiv:2310.04948*.
- Chen, X., Soh, B. W., Ooi, Z. E., et al. (2024). Constructing custom thermodynamics using deep learning. *Nature Computational Science*, 4:66–85.
- Cheng, Q., Liu, L., Zhu, Q., Yu, R., Jin, Z., Xie, Y., and Jia, X. (2025). Llm-based evaluation policy extraction for ecological modeling.
- Chiuso, A. and Pillonetto, G. (2019). System identification: A machine learning perspective. *Annual Review of Control, Robotics, and Autonomous Systems*, 2(1):281–304.

- Condie, S. A., Anthony, K. R. N., Babcock, R. C., Baird, M. E., Beeden, R., Fletcher, C. S., Gorton, R., Harrison, D., Hobday, A. J., Plagányi, E. E., and Westcott, D. A. (2021). Large-scale interventions may delay decline of the Great Barrier Reef. *Royal Society Open Science*, 8:201296.
- Dichmont, C. M., Deng, R. A., Punt, A. E., Brodziak, J., Chang, Y.-J., Cope, J. M., Ianelli, J. N., Legault, C. M., Methot Jr, R. D., Porch, C. E., et al. (2017). From data rich to data-limited harvest strategies—does more data mean better management? *ICES Journal of Marine Science*, 74(3):670–686.
- Edwards, A. M. and Brindley, J. (1999). Zooplankton mortality and the dynamical behaviour of plankton population models. *Bulletin of mathematical biology*, 61(2):303–339.
- Ellner, S. P., Seifu, Y., and Smith, R. H. (2002). Fitting population dynamic models to time-series data by gradient matching. *Ecology*, 83(8):2256–2270.
- Floryan, D. and Graham, M. D. (2022). Data-driven discovery of intrinsic dynamics. *Nature Machine Intelligence*, 4:1113–1120.
- Franks, P. J. (2002). Npz models of plankton dynamics: their construction, coupling to physics, and application. *Journal of Oceanography*, 58(2):379–387.
- Gandhi, Y., Zheng, K., Jha, B., Nomura, K.-i., Nakano, A., Vashishta, P., and Kalia, R. K. (2024). Generative AI-driven forecasting of oil production. *arXiv preprint arXiv:2409.16482*.
- Gårdmark, A., Lindegren, M., Neuenfeldt, S., Blenckner, T., Heikinheimo, O., Müller-Karulis, B., Niiranen, S., Tomczak, M. T., Aro, E., Wikström, A., et al. (2013). Biological ensemble modeling to evaluate potential futures of living marine resources. *Ecological applications*, 23(4):742–754.
- Garza, A., Challu, C., and Mergenthaler-Canseco, M. (2023). Timegpt-1. *arXiv preprint arXiv:2310.03589*.
- Gauthier, P. (2024). Aider: AI pair programming in your terminal.
- Gougherty, A. V. and Clipp, H. L. (2024). Testing the reliability of an ai-based large language model to extract ecological information from the scientific literature. *npj Biodiversity*, 3(1):13.

- Hassani, H. and Silva, E. S. (2024). Predictions from Generative Artificial Intelligence Models: Towards a New Benchmark in Forecasting Practice. *Information*, 15(6):291.
- Holden, M. H., Plagányi, E. E., Fulton, E. A., Campbell, A. B., Janes, R., Lovett, R. A., Wickens, M., Adams, M. P., Botelho, L. L., Dichmont, C. M., Erm, P., Helmstedt, K. J., Heneghan, R. F., Mendiolar, M., Richardson, A. J., Rogers, J. G. D., Saunders, K., and Timms, L. (2024). Cost–benefit analysis of ecosystem modeling to support fisheries management. *Journal of Fish Biology*, 104(6):1667–1674.
- Holden, P. B. and Ellner, S. P. (2016). Prediction in ecology: a first-principles framework. *Ecological Monographs*, 86(3):303–319.
- Huntingford, C., Nicoll, A. J., Klein, C., and Ahmad, J. A. (2024). Potential for Equation Discovery with AI in the Climate Sciences.
- Karniadakis, G. E., Kevrekidis, I. G., Lu, L., Perdikaris, P., Wang, S., and Yang, L. (2021). Physics-informed machine learning. *Nature Reviews Physics*, 3:422–440.
- Keck, F., Broadbent, H., and Altermatt, F. (2025). Extracting massive ecological data on state and interactions of species using large language models. *bioRxiv*, pages 2025–01.
- Kramer, S., Cerrato, M., Džeroski, S., and King, R. (2023). Automated scientific discovery: from equation discovery to autonomous discovery systems. *arXiv preprint arXiv:2305.02251*.
- Kristensen, K. (2014). TMB: General random effect model builder tool inspired by ADMB. *R package version*, 1.
- Li, H., Liu, J., Wang, Z., Luo, S., Jia, X., and Yao, H. (2024). LITE: Modeling Environmental Ecosystems with Multimodal Large Language Models. *arXiv preprint arXiv:2404.01165*.
- Malhi, Y., Franklin, J., Seddon, N., Solan, M., Turner, M. G., Field, C. B., and Knowlton, N. (2020). Climate change and ecosystems: threats, opportunities and solutions.
- Mammides, C. and Papadopoulos, H. (2024). The role of large language models in interdisciplinary research: Opportunities, challenges and ways forward. *Methods in Ecology and Evolution*, 15(10):1774–1776.

- McCarthy, M. A. and Possingham, H. P. (2004). Active adaptive management for conservation. *Conservation Biology*, 18(2):534–541.
- Morales-García, J., Llanes, A., Arcas-Túnez, F., and Terroso-Sáenz, F. (2024). Developing time series forecasting models with generative large language models. *ACM Transactions on Intelligent Systems and Technology*.
- Morello, E. B., Plagányi, E. E., Babcock, R. C., Sweatman, H., Hillary, R., and Punt, A. E. (2014). Model to manage and reduce crown-of-thorns starfish outbreaks. *Marine Ecology Progress Series*, 512:167–183.
- Munch, S. B., Rogers, T. L., and Sugihara, G. (2023). Recent developments in empirical dynamic modelling. *Methods in Ecology and Evolution*, 14(3):732–745.
- Patterson, K., Cook, R., Darby, C., Gavaris, S., Kell, L., Lewy, P., Mesnil, B., Punt, A., Restrepo, V., Skagen, D. W., et al. (2001). Estimating uncertainty in fish stock assessment and forecasting. *Fish and fisheries*, 2(2):125–157.
- Plagányi, É. E., Babcock, R. C., Rogers, J., Bonin, M., and Morello, E. B. (2020). Ecological analyses to inform management targets for the culling of crown-of-thorns starfish to prevent coral decline. *Coral Reefs*, 39(5):1483–1499.
- Plagányi, E. E., Punt, A. E., Hillary, R., Morello, E. B., Thébaud, O., Hutton, T., Pillans, R. D., Thorson, J. T., Fulton, E. A., Smith, A. D. M., Smith, F., Bayliss, P., Haywood, M., Lyne, V., and Rothlisberg, P. C. (2014). Multispecies fisheries management and conservation: tactical applications using models of intermediate complexity. *Fish and Fisheries*, 15:1–22.
- Rogers, J. G., Plagányi, É. E., Babcock, R. C., Fletcher, C. S., and Westcott, D. A. (2023). Improving coral cover using an integrated pest management framework. *Ecological Applications*, 33(8):e2913.
- Rogers, J. G., Plagányi, É. E., Blamey, L. K., and Desbiens, A. A. (2024). Validating effectiveness of crown-of-thorns starfish control thresholds to limit coral loss throughout the great barrier reef. *Coral Reefs*, pages 1–16.
- Rogers, J. G. D. and Plagányi, E. E. (2022). Culling corallivores improves short-term coral recovery under bleaching scenarios. *Nature Communications*, 13:2520.

- Schaeffer, H. (2017). Learning partial differential equations via data discovery and sparse optimization. *Proceedings of the Royal Society A*, 473:20160446.
- Serper (2024). Serper search api.
- Spillias, S., Adams, M., Boschetti, F., Constable, A., Dunstan, P., Ferrier, S., Porobic, J., Gorton, B., Grimberg, E., Grigg, N., Harfoot, M., Hirsch, P., Hobday, A. J., Holden, M., Hutton, T., Kaur, S., Melbourne-Thomas, J., Paris, C., Parker, D., Plagányi-Lloyd, E., Rogers, J. G. D., Stitzlein, C., Trebilco, R., Weiley, V., Wild-Allen, K., Woolley, S. N. C., and Fulton, E. A. (2024a). The Future of Artificial Intelligence in Ecosystem Modelling. *Manuscript in preparation*.
- Spillias, S., Ollerhead, K., Andreotta, M., Annand-Jones, R., Boschetti, F., Duggan, J., Karcher, D., Paris, C., Shellock, R., and Trebilco, R. (2024b). Evaluating generative ai to extract qualitative data from peer-reviewed documents.
- Su, J., Jiang, C., Jin, X., Qiao, Y., Xiao, T., Ma, H., Wei, R., Jing, Z., Xu, J., and Lin, J. (2024). Large language models for forecasting and anomaly detection: A systematic literature review. *arXiv preprint arXiv:2402.10350*.
- Vollert, S. A., Drovandi, C., and Adams, M. P. (2024). Unlocking ensemble ecosystem modelling for large and complex networks. *PLoS Computational Biology*, 20(3):e1011976.
- Wang, Z., Zhang, Y., Zhao, Y., Chen, Y., and Li, P. (2024). BioResearcher: Automating Biomedical Research via Large Language Models. *arXiv preprint arXiv:2412.09429*.
- Weiskopf, S. R., Rubenstein, M. A., Crozier, L. G., Gaichas, S., Griffis, R., Halofsky, J. E., Hyde, K. J., Morelli, T. L., Morisette, J. T., Muñoz, R. C., et al. (2020). Climate change effects on biodiversity, ecosystems, ecosystem services, and natural resource management in the united states. *Science of the Total Environment*, 733:137782.
- Willard, J., Jia, X., Xu, S., Steinbach, M., and Kumar, V. (2022). Integrating scientific knowledge with machine learning for engineering and environmental systems. *ACM Computing Surveys*, 55(4):1–37.
- Wills, S., Poon, S. T., Salili-James, A., and Scott, B. (2024). The use of generative ai for coding in academia. *Methods in Ecology and Evolution*, 15(12):2189–2191.

Xu, Y., Liu, X., Cao, X., Huang, C., Liu, E., Qian, S., Liu, X., Wu, Y., Dong, F., and Qiu, C.-W. (2021). Artificial intelligence: A powerful paradigm for scientific research. *The Innovation*, 2(2):100179.

Zhang, X., Chowdhury, R. R., Gupta, R. K., and Shang, J. (2024). Large language models for time series: A survey. *arXiv preprint arXiv:2402.01801*.